



Agave v4.2

Security Assessment

Prepared for
Anza

Stepan Fedotov
Gabriel Schneider
Thiago Tavares
Zelda Hodel

Prepared by
Otter Audits, LLC

Prepared on
August 19th, 2026

trixter@osec.io
gbrls@osec.io
thitav@osec.io
zel@osec.io

Table of Contents

1. Executive Summary	2
1.1. Overview	2
1.2. Key Findings	2
2. Scope	3
2.1. Transaction Execution (SVM/Runtime)	3
2.2. Staking, Rewards & Commission	3
2.3. Transaction Scheduling & Banking	3
2.4. Program Loading & Execution (sBPF)	3
2.5. Consensus & Voting (Tower BFT)	4
A. Vulnerability rating scale	5
B. Procedure	6

1. Executive Summary

1.1. Overview

Anza engaged OtterSec to conduct a security assessment of the agave v4.2 release (v4.1.0..v4.2.0-beta.0). This assessment was conducted between July 6th and August 1st, 2026. For information on our auditing methodology, refer to [Appendix B](#).

1.2. Key Findings

No critical or high findings were produced throughout this audit engagement.

2. Scope

The source code was delivered to us in a Git repository at <https://github.com/anza-xyz>.

This audit was performed against the following PRs:

2.1. Transaction Execution (SVM/Runtime)

[PR#10127](#), [PR#11814](#), [PR#11880](#), [PR#12183](#), [PR#12264](#), [PR#12284](#), [PR#12384](#), [PR#12406](#), [PR#12413](#),
[PR#12440](#), [PR#12444](#), [PR#12448](#), [PR#12490](#), [PR#12502](#), [PR#12503](#), [PR#12519](#), [PR#12538](#), [PR#12540](#),
[PR#12548](#), [PR#12561](#), [PR#12589](#), [PR#12590](#), [PR#12620](#), [PR#12646](#), [PR#12659](#), [PR#12683](#), [PR#12723](#),
[PR#12733](#), [PR#12738](#), [PR#12743](#), [PR#12745](#), [PR#12756](#), [PR#12795](#), [PR#12814](#), [PR#12820](#), [PR#12834](#),
[PR#12838](#), [PR#12846](#), [PR#12857](#), [PR#12874](#), [PR#12880](#), [PR#12881](#), [PR#12886](#), [PR#12890](#), [PR#12891](#),
[PR#12898](#), [PR#12906](#), [PR#12907](#), [PR#12913](#), [PR#12919](#), [PR#12921](#), [PR#12940](#), [PR#12941](#), [PR#12990](#),
[PR#13037](#), [PR#13060](#), [PR#13074](#), [PR#13079](#), [PR#13093](#), [PR#13096](#), [PR#13120](#), [PR#13123](#), [PR#13125](#),
[PR#13152](#), [PR#13168](#), [PR#13187](#), [PR#13203](#), [PR#13204](#), [PR#13233](#), [PR#13258](#), [PR#13340](#), [PR#13341](#),
[PR#13346](#), [PR#13347](#), [PR#13353](#), [PR#13364](#), [PR#13402](#), [PR#13403](#), [PR#13404](#), [PR#13405](#), [PR#13406](#),
[PR#13417](#), [PR#13427](#), [PR#13435](#), [PR#13439](#), [PR#13446](#), [PR#13467](#), [PR#13479](#).

2.2. Staking, Rewards & Commission

[PR#9637](#), [PR#12096](#), [PR#12317](#), [PR#12345](#), [PR#12408](#), [PR#12539](#), [PR#12550](#), [PR#12588](#), [PR#12815](#),
[PR#12872](#), [PR#12966](#), [PR#13051](#), [PR#13061](#), [PR#13153](#), [PR#13229](#), [PR#13257](#), [PR#13285](#), [PR#13295](#),
[PR#13355](#), [PR#13373](#), [PR#13532](#).

2.3. Transaction Scheduling & Banking

[PR#12263](#), [PR#12422](#), [PR#12485](#), [PR#12739](#), [PR#12746](#), [PR#12773](#), [PR#12779](#), [PR#12902](#), [PR#12954](#),
[PR#12963](#), [PR#12984](#), [PR#13027](#), [PR#13092](#), [PR#13283](#), [PR#13291](#).

2.4. Program Loading & Execution (sBPF)

[PR#9663](#), [PR#12134](#), [PR#12382](#), [PR#12407](#), [PR#12410](#), [PR#12423](#), [PR#12461](#), [PR#12465](#), [PR#12475](#),
[PR#12484](#), [PR#12498](#), [PR#12518](#), [PR#12520](#), [PR#12521](#), [PR#12523](#), [PR#12577](#), [PR#12579](#), [PR#12591](#),
[PR#12592](#), [PR#12603](#), [PR#12605](#), [PR#12607](#), [PR#12611](#), [PR#12612](#), [PR#12619](#), [PR#12622](#), [PR#12623](#),
[PR#12630](#), [PR#12644](#), [PR#12658](#), [PR#12675](#), [PR#12677](#), [PR#12684](#), [PR#12692](#), [PR#12709](#), [PR#12710](#),
[PR#12717](#), [PR#12725](#), [PR#12727](#), [PR#12735](#), [PR#12747](#), [PR#12748](#), [PR#12749](#), [PR#12788](#), [PR#12792](#),
[PR#12793](#), [PR#12804](#), [PR#12938](#), [PR#12948](#), [PR#12960](#), [PR#12977](#), [PR#13011](#), [PR#13012](#), [PR#13017](#),
[PR#13019](#), [PR#13045](#), [PR#13046](#), [PR#13069](#), [PR#13070](#), [PR#13072](#), [PR#13078](#), [PR#13081](#), [PR#13095](#),
[PR#13104](#), [PR#13108](#), [PR#13126](#), [PR#13128](#), [PR#13131](#), [PR#13154](#), [PR#13156](#), [PR#13206](#), [PR#13217](#),
[PR#13219](#), [PR#13223](#), [PR#13224](#), [PR#13247](#), [PR#13264](#), [PR#13265](#), [PR#13266](#), [PR#13274](#), [PR#13276](#),
[PR#13277](#), [PR#13298](#), [PR#13301](#), [PR#13302](#), [PR#13313](#), [PR#13323](#), [PR#13324](#), [PR#13336](#), [PR#13345](#).

[PR#13348](#), [PR#13399](#), [PR#13440](#), [PR#13442](#), [PR#13444](#), [PR#13445](#), [PR#13465](#), [PR#13473](#), [PR#13478](#),
[PR#13484](#), [PR#13486](#), [PR#13487](#), [PR#13489](#), [PR#13527](#), [PR#13528](#), [PR#13656](#).

25. Consensus & Voting (Tower BFT)

[PR#12176](#), [PR#13362](#), [PR#13367](#), [PR#13372](#), [PR#13394](#).

Appendix A.

Vulnerability rating scale

We rated our findings according to the following scale. Vulnerabilities have immediate security implications. Informational findings may be found in the general findings.

Severity	Description
● Critical	Vulnerabilities that immediately result in a loss of availability, consensus divergence, or exploitable memory corruption with minimal preconditions.
● High	Vulnerabilities that may result in consensus divergence or denial of service but require specific transaction or network conditions to exploit.
● Medium	Vulnerabilities that may result in incorrect behavior or state divergence under specific configurations or edge cases.
● Low	Low probability vulnerabilities, which are still exploitable but require extenuating circumstances or undue risk.
● Informational	Best practices to mitigate future security risks. These are classified as general findings.

Resolution status is one of  Resolved,  Ack'd, or  To-do.

Appendix B.

Procedure

As part of our standard auditing procedure, we split our analysis into two main sections: design and implementation.

The design review asks whether the programs's architecture and operating assumptions support its intended behavior. We study the inputs it accepts, the decisions it makes, the information shared between components, and the services or participants it trusts. We then consider how it behaves when those inputs are unexpected or conflicting. The goal is to find conditions that can produce an incorrect decision, inconsistent state, or loss of service.

Design problems often arise from interactions rather than one defective function. Two components may each process a value correctly but interpret it differently. A failure may also leave one component ready to continue while another treats the same operation as incomplete. Local checks can look correct in both cases, so we follow the complete operation across component boundaries.

The implementation review asks whether the code matches the design. We trace untrusted data from its entry point to the point where it affects behavior, checking the assumptions made along the way. The review covers validation, error handling, state updates, concurrency, persistence, resource use, and cleanup. Its exact focus depends on the target's language, runtime, and architecture rather than a fixed list of bug classes.

The boundary between design and implementation is not exact. Missing validation may be a local coding error, while the same omission across several components may reflect an unclear rule. We classify a finding by its root cause, but judge its impact from the validator's resulting behavior and the conditions needed to trigger it.

We begin by reading the target as a team and agreeing on its main components, trust boundaries, and expected behavior. Auditors then review different paths through the system and compare notes where those paths meet. We reproduce suspected issues with focused tests or traces when the setup allows. Before reporting a finding, another auditor reviews its trigger, impact, and proposed remediation.

While sometimes the line between design and implementation can be blurry, we hope this gives some insight into our auditing procedure and thought process